

Transfer Learning Assignment

Download all the data in this [rar_file](#), it contains all the data required for the assignment. When you unrar the file you'll get the files in the following format:
path/to/the/image.tif.category

where the categories are numbered 0 to 15, in the following order:

```
0 letter
1 form
2 email
3 handwritten
4 advertisement
5 scientific report
6 scientific publication
7 specification
8 file folder
9 news article
10 budget
11 invoice
12 presentation
13 questionnaire
14 resume
15 memo
```

There is a file named as 'labels_final.csv', it consists of two columns. First column is path which is the required path to the images and second is the class label.

```
In [ ]: #the dataset that you are dealing with is quite large 3.7 GB and hence there are two methods to import the data to Colab
# Method 1- you can use gdown module to get the data directly from Google drive to Colab
# the syntax is as follows !gdown --id file_id , for ex - running the below cell will import the rvl-cdip.rar dataset
```

```
In [ ]: !gdown --id 1Z4TyI7FcFVEEx8qd14j09qvxqaL5qoEu
```

Access denied with the following error:

Cannot retrieve the public link of the file. You may need to change the permission to 'Anyone with the link', or have had many accesses.

You may still be able to access the file from the browser:

<https://drive.google.com/uc?id=1Z4TyI7FcFVEEx8qd14j09qvxqaL5qoEu>

```
In [ ]: # Method -2 you can also import the data using wget function
#https://www.youtube.com/watch?v=BPUpfVq7RaY
```

```
In [ ]: !curl --header "Host: doc-0s-80-docs.googleusercontent.com" --header "User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/98.0.4758.102 Safari/537.36" --header "Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header "Accept-Language: en-US,en;q=0.9" --header "Referer: https://drive.google.com/" --header "Cookie: AUTH_nf14kr63csq6es11jms1j175edsrf0ud_nonce=hpi3n12752t54; ab.storage.deviceId.7af503ae-0c84-478f-98b0-ecffff5d67750=%7B%22g%22%3A%22cf6f5c5c-eb07-f7a6-15ec-eeg63d656f051%22%2C%22c%22%3A1641395211601%22%22%3A1645860762542%22%22%3A1645860762563%7D; ab.storage.userId.7af503ae-0c84-478f-98b0-ecffff5d67750=%7B%22g%22%3A%22334488%22%2C%22c%22%3A1645860762542%22%22%3A1645860768738%2C%22c%22%3A1645860768738%7D" --header "Connection: keep-alive" "https://doc-0s-80-docs.googleusercontent.com/docs/securesc/crnfs1fh7cm9739i4jrpvn7r10e8dd/16ul0t7nnh6911ekugj2h2arggrpmbho/1646206425000/00484516897554883881/15475546427455642759/1247yI7FcFVEEx8qd14j09qvxqaL5qoEu?ax=ACxEAsb9P017nsrkXNZE5Pej_vgnVnD2gpPS10UPxAcP6bluiUryYiuFezkEmYdQ3af3yUfeB1mKoz-3QzQ2G2FfyfGhaal6JJJEh1BGWlH9rfkx7MvsCZggRGTfjuoWo2sEX9ka2pGjhCDpqxSgSVoh7s4cZGnfGyVraJuztRKzGjtV7q6LMDSpq83H8K7H1fY_Kp6s4Ne-lpWkewZQEEbf-hf23leJn5U7eognZl1HQps8hewmUDbtnjKLOX02Xmh4IkY1FXT8FHJuDqvJGZJ91-3p2eVehD6BqcujoTuAygbs56n8t7wshTp50jAepTxUVRLi00gWAlLo7it1NzY8_90m67dokkk-F1fkawRJ6NLy9qgf18bLP_LLDFQYWGmybwpyTayn_vapuaqIumBvQsloZrt8LuhZ7qTOHFQfbXTIZQutccAK101EQP60XI999I30XThzL7Eg6f0C9qYEsnYFuapxF7R9ruRz_3Ys1wrpNXAJ0tFenBrAIGeuOmUhizIKst0MB_-jFc9vbEt2CpnXix9ZppU9xZdwR13sWPxrKP_JEasRwP08MtvrPj1aDB-67-25HcoSYPTXCJVDKEduvukGVRc0dCguRrGVehgpgLttCckid5qsPmYkwMv-n1MoYtsr2ZAw7-7HYOPM&authuser=0&nonce=hpi3n12752t54&user=15475546427455647259&hash=4v1rf1mg7rgc55m65vg8f3fbm4bnvll1" -L -o "rv1-cdip.rar"
```

% Total	% Received	% Xferd	Average Speed	Time Dload	Time Upload	Time Total	Time Spent	Time Left	Current Speed
0	0	0	0	0	0	0	--::--	--::--	--::-- 0

```
In [ ]: #unrar the file
get_ipython().system_raw("unrar x rvl-cdip.rar")
```

2. On this image data, you have to train 3 types of models as given below You have to split the data into Train and Validation data.

1. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data. or you can use this method also <https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

Note- In the reference notebook you were dealing with jpg images, in the given dataset you are dealing with tiff images. Imagedatagenerator works with both type of images. If you want to use custom data pipeline then you have to convert your tiff images to jpg images.

1. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.
2. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)
1. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

1. Do print model.summary() and draw model_plots for each of the model.

Model-1

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and an output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer
4. Print model.summary() and plot the architecture of the model.
[Reference for plotting model](#)
5. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

In []:

In []:

Model-2

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. Any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size $7 \times 7 \times 512$ can be equivalently expressed as a CONV layer with $F=7, P=0, S=1, K=4096$. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be $1 \times 1 \times 4096$ since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](#) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer
4. Print model.summary() and plot the architecture of the model.
[Reference for plotting model](#)
5. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

In []:

In []:

Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

In []:

In []:

Please write your observations or a brief summary of the results that you get after performing transfer learning with reference to model1, model2 and model3

In []:

In []:

```

from keras.optimizers import RMSprop
from keras.callbacks import ModelCheckpoint, Callback, EarlyStopping
from keras.optimizers import Adam
from keras.utils import np_utils
from keras.models import Sequential
from keras.layers import Input, Dropout, Flatten, Conv2D, MaxPooling2D, Dense, Activation
from keras.optimizers import RMSprop
from keras.callbacks import ModelCheckpoint, Callback, EarlyStopping
from keras.utils import np_utils
from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img
#!pip install opencv_python
import cv2

```

```

In [ ]: !curl --header "Host: doc-0s-80-docs.googleusercontent.com" --header "User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/98.0.4758.82 Safari/537.36" --header "Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header "Accept-Language: en-US,en;q=0.9" --header "Referer: https://drive.google.com" --header "Cookie: AUTH_nf14kr63csq6es11jms1j175edsrf0ud_nonce=gp0sj6pjcqf6t0" --header "Connection: keep-alive" "https://doc-0s-80-docs.googleusercontent.com/docs/securesc/crnfsi1hu7cm973914jrpvn7r1008dd/iqqs1vme4p5msrebboeqif31ufgou18q/1646342625000/00484516897554883881/15475546427455647259/1247y17FcFVEx8qd14j09qxvxqaLsQoE?ax=ACxExAsalwU-gc9wvBrEv1Q0FmV2fPt5v67H3BXrsv_vDXd05CTyToatqhSudyCkk_4MMWF880PquENPglk3sWikh1V0qVBuk1bOvequ9Eaq693-jjiG1P-5DX-pcZUoQyr9c1XXnyL-X3h3TltBeE9D_SqNFBSs_Whx1l0cZz7vzTGIxm_2K-E_Yqw1ffcORNIrLkoJTV1caaCRSugDggnp5CobHL7jMKH7e1nTbQIUhahTkSmA_VUOSNg3xObbWlBARq7G6-syAYpnucPEkfsgseHKKR9kccdVGNHScnuctrvF1vFUUDfvJ5P2Ywx-J3jY5wYcxOUxMu0F7A7sxv1uhSRJK89Gea3pjK-pNBjh9bGxUuswOUA0Xuu0N1G6SbhHDVGafIA18qP8je5rUtwmVHuVyyugpQp15rNl_0nWnQplii_0pZLxWMrDoyn5hClnJfL5fydH2bMlw2CStGvooV3Yw3A0QR3qKibG-JJFBBvh1cUTV24Y1fr8NMaTfphtKuriwpwqqxfgd8tU6c1h2X8Gaij3YhH4A8lTupCmwbj8lMqRZDQkiyQubf2QbkWLks5k1eG440KEX8ujrFwIZU18s1m-pj80dUqLcxGHYc5fXIqKxv17Bz22pEJ7NPzHTVsS_xv3V7thMXctjsUCHg&authuser=0&nonce=gp0sj6pjcq6t0&user=15475546427455647259&hash=824dscvm6fos4k1f8vlef3nm2kekuek" -L -o "rvl-cdip.rar"

```

% Total	% Received	% Xferd	Average Speed	Time Dload	Time Upload	Time Total	Time Spent	Time Left	Current Speed
100	4444M	100	4444M	0	0	175M	0	0:00:25	0:00:25 ---:-- 177M

```

In [ ]: #unrar the file
get_ipython().system_raw("unrar x rvl-cdip.rar")

```

```

In [ ]: #import all the required Libraries
import tensorflow as tf
import os
import numpy as np
import pandas as pd

```

```

df=pd.read_csv('labels_final.csv',dtype=str)

```

```

In [ ]: df.head(2)

```

```

Out[ ]:
path    label
0   imagesv/v/o/h/voh71d00/509132755+-2755.tif      3
1   imagesv/l/x/t/xt19d00/502213303.tif      3

```

```

In [ ]: #new model

```

```

In [ ]: import numpy as np
import pandas as pd
import keras
import pydotplus
from sklearn.tree import export_graphviz

```

```

In [ ]: from keras.utils.vis_utils import plot_model

```

```

In [ ]: images_df = pd.read_csv('labels_final.csv')
images_df['label'] = images_df['label'].astype(str)
images_dir = 'data_final'

```

```

In [ ]: import tensorflow as tf
target_size = (128, 128)
data_generator=tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255,validation_split=0.2)
train_images_generator = data_generator.flow_from_dataframe(images_df, images_dir,x_col='path',y_col='label',class_mode='categorical', target_size=target_size,batch_size=128, subset='training')
valid_images_generator = data_generator.flow_from_dataframe(images_df, images_dir,x_col='path',y_col='label',class_mode='categorical', target_size=target_size,batch_size=128, subset='validation')
#images_df act is csv file which store files of location image_dir.
# i was not knowing we can use csv file anf

```

Found 38400 validated image filenames belonging to 16 classes.
Found 9600 validated image filenames belonging to 16 classes.

```

In [ ]: vgg16base_model = tf.keras.applications.VGG16(weights='imagenet', include_top=False,input_shape=(*target_size, 3))
for layer in vgg16base_model.layers:
    layer.trainable = False
conv_layer = tf.keras.layers.Conv2D(512, (3, 3), activation='relu',name='Conv')(vgg16base_model.output)
max_pool_layer = tf.keras.layers.MaxPooling2D(name='Max_Pool')(conv_layer)
flatten = tf.keras.layers.Flatten(name='Flatten')(max_pool_layer)
Fully_Connected_1 = tf.keras.layers.Dense(1024,activation='relu',name='Fully_Connected_Layer1')(flatten)
Fully_Connected_2 = tf.keras.layers.Dense(256,activation='relu',name='Fully_Connected_Layer2')(Fully_Connected_1)
Output = tf.keras.layers.Dense(16, activation='softmax',name='Output')(Fully_Connected_2)
model = tf.keras.models.Model(inputs=vgg16base_model.inputs, outputs=Output)
model.compile(optimizer='Adam', loss='categorical_crossentropy',metrics=['accuracy'])

```

```

In [ ]: model.summary()

```

Model: "model_1"

Layer (type)	Output Shape	Param #
<hr/>		
input_2 (InputLayer)	[(None, 128, 128, 3)]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
<hr/>		

dm
num

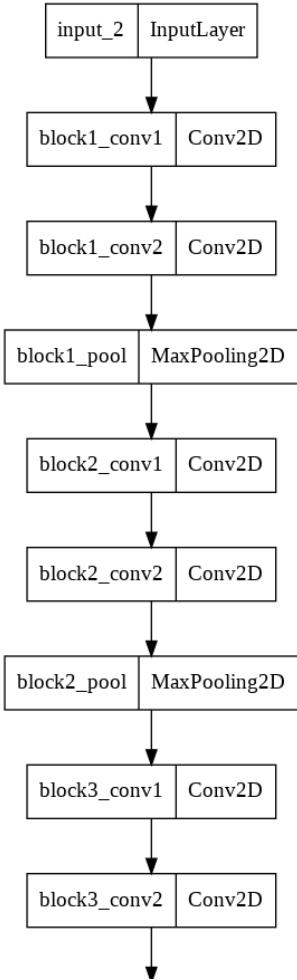
```

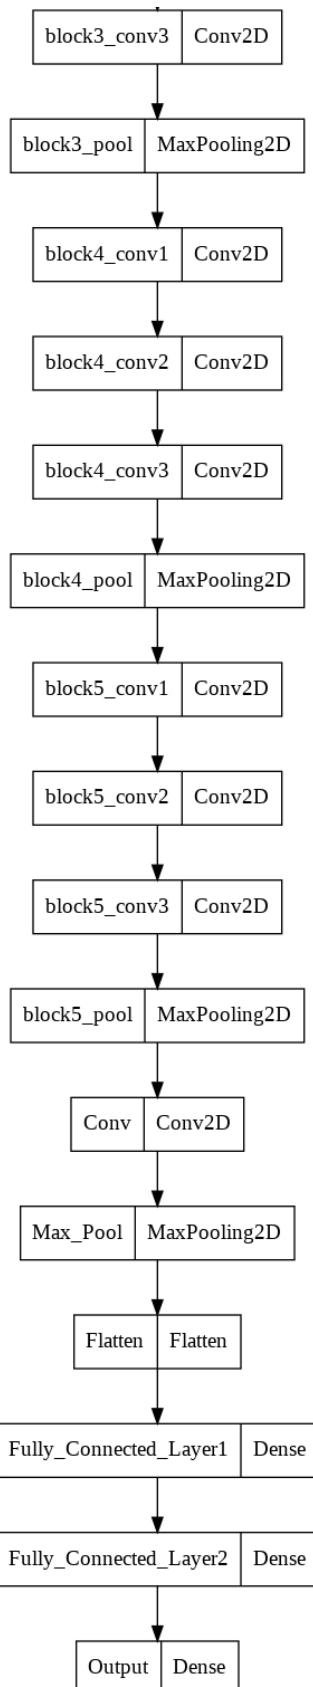
block1_pool (MaxPooling2D) (None, 64, 64, 64) 0
block2_conv1 (Conv2D) (None, 64, 64, 128) 73856
block2_conv2 (Conv2D) (None, 64, 64, 128) 147584
block2_pool (MaxPooling2D) (None, 32, 32, 128) 0
block3_conv1 (Conv2D) (None, 32, 32, 256) 295168
block3_conv2 (Conv2D) (None, 32, 32, 256) 590080
block3_conv3 (Conv2D) (None, 32, 32, 256) 590080
block3_pool (MaxPooling2D) (None, 16, 16, 256) 0
block4_conv1 (Conv2D) (None, 16, 16, 512) 1180160
block4_conv2 (Conv2D) (None, 16, 16, 512) 2359808
block4_conv3 (Conv2D) (None, 16, 16, 512) 2359808
block4_pool (MaxPooling2D) (None, 8, 8, 512) 0
block5_conv1 (Conv2D) (None, 8, 8, 512) 2359808
block5_conv2 (Conv2D) (None, 8, 8, 512) 2359808
block5_conv3 (Conv2D) (None, 8, 8, 512) 2359808
block5_pool (MaxPooling2D) (None, 4, 4, 512) 0
Conv (Conv2D) (None, 2, 2, 512) 2359808
Max_Pool (MaxPooling2D) (None, 1, 1, 512) 0
Flatten (Flatten) (None, 512) 0
Fully_Connected_Layer1 (Dense) (None, 1024) 525312
Fully_Connected_Layer2 (Dense) (None, 256) 262400
Output (Dense) (None, 16) 4112
=====
Total params: 17,866,320
Trainable params: 3,151,632
Non-trainable params: 14,714,688

```

In []: `plot_model(model)# use to describe model`

Out[]:





```

In [ ]: import tensorflow
batch_size = 500
steps = len(train_images_generator.labels) // batch_size
early_stopping = keras.callbacks.EarlyStopping(patience=2)
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard= tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,write_graph=True,write_grads=True)
callbacks_list = [early_stopping,tensorboard]
history = model.fit(train_images_generator,validation_data=valid_images_generator,epochs=3,steps_per_epoch=steps,callbacks=callbacks_list)

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
Epoch 1/3
76/76 [=====] - 3286s 43s/step - loss: 1.9560 - accuracy: 0.3809 - val_loss: 1.5800 - val_accuracy: 0.
5140
Epoch 2/3
  
```

```
76/76 [=====] - 3290s 43s/step - loss: 1.4735 - accuracy: 0.5427 - val_loss: 1.4426 - val_accuracy: 0.  
5483  
Epoch 3/3  
76/76 [=====] - 3307s 44s/step - loss: 1.3273 - accuracy: 0.5892 - val_loss: 1.2983 - val_accuracy: 0.  
6032
```

```
In [ ]: # Save the weights using the `checkpoint_path` format  
checkpoint_path = "/content/"  
# Save the weights  
model.save_weights('./checkpoints/my_checkpoint')
```

```
In [ ]: # Restore the weights  
model.load_weights('./checkpoints/my_checkpoint')
```

```
Out[ ]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f8e29bee0d0>
```

```
In [ ]: %reload_ext tensorboard  
%tensorboard --logdir {log_dir}
```

```
In [ ]: ##### model-2
```

```
In [ ]: vgg16base_model2 = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape=(*target_size, 3))  
for layer in vgg16base_model2.layers:  
    layer.trainable = False  
conv_layer1 = tf.keras.layers.Conv2D(64, 4, 4, padding='same', activation='relu', name='Conv1')(vgg16base_model2.output)  
#conv_layer1 = tf.keras.layers.Conv2D(512, kernel_size=(4, 4), padding='valid') ,4,4 BROUGHT O/P OF 1,1,64  
conv_layer2 = tf.keras.layers.Conv2D(128, (3, 3), activation='relu', padding='same', name='Conv2')(conv_layer1)  
flatten = tf.keras.layers.Flatten(name='Flatten')(conv_layer2)  
output = tf.keras.layers.Dense(16, activation='softmax', name='Output')(flatten)  
model2 = tf.keras.models.Model(inputs=vgg16base_model2.inputs, outputs=output, name='Model_2')  
model2.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

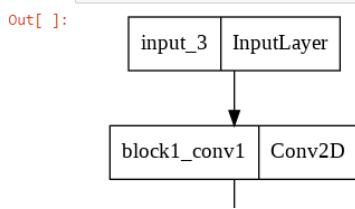
```
In [ ]: model2.summary()
```

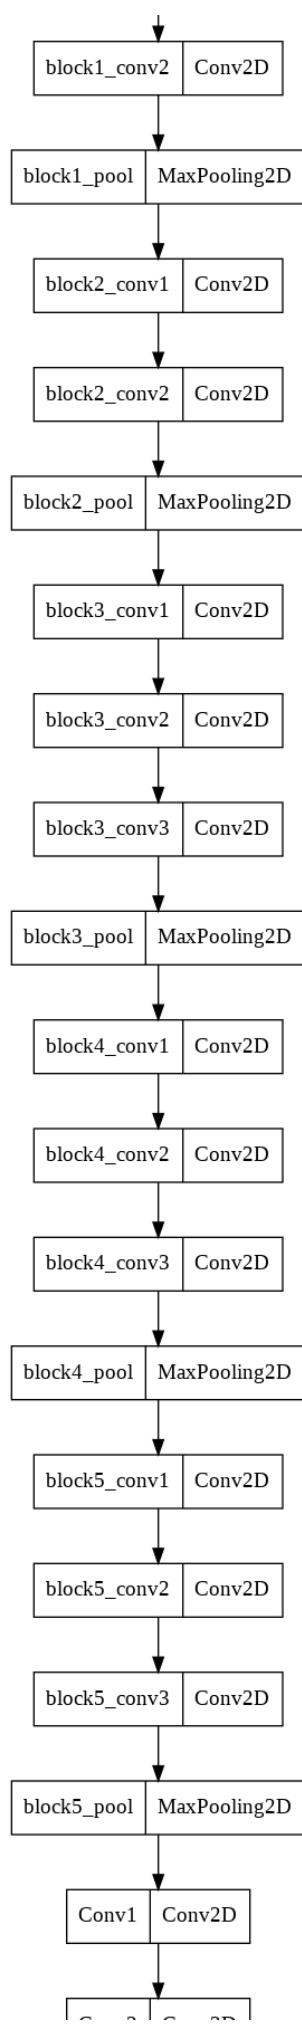
```
Model: "Model_2"  

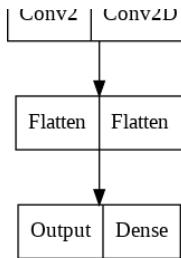

| Layer (type)               | Output Shape         | Param # |
|----------------------------|----------------------|---------|
| input_3 (InputLayer)       | [None, 128, 128, 3]  | 0       |
| block1_conv1 (Conv2D)      | (None, 128, 128, 64) | 1792    |
| block1_conv2 (Conv2D)      | (None, 128, 128, 64) | 36928   |
| block1_pool (MaxPooling2D) | (None, 64, 64, 64)   | 0       |
| block2_conv1 (Conv2D)      | (None, 64, 64, 128)  | 73856   |
| block2_conv2 (Conv2D)      | (None, 64, 64, 128)  | 147584  |
| block2_pool (MaxPooling2D) | (None, 32, 32, 128)  | 0       |
| block3_conv1 (Conv2D)      | (None, 32, 32, 256)  | 295168  |
| block3_conv2 (Conv2D)      | (None, 32, 32, 256)  | 590080  |
| block3_conv3 (Conv2D)      | (None, 32, 32, 256)  | 590080  |
| block3_pool (MaxPooling2D) | (None, 16, 16, 256)  | 0       |
| block4_conv1 (Conv2D)      | (None, 16, 16, 512)  | 1180160 |
| block4_conv2 (Conv2D)      | (None, 16, 16, 512)  | 2359808 |
| block4_conv3 (Conv2D)      | (None, 16, 16, 512)  | 2359808 |
| block4_pool (MaxPooling2D) | (None, 8, 8, 512)    | 0       |
| block5_conv1 (Conv2D)      | (None, 8, 8, 512)    | 2359808 |
| block5_conv2 (Conv2D)      | (None, 8, 8, 512)    | 2359808 |
| block5_conv3 (Conv2D)      | (None, 8, 8, 512)    | 2359808 |
| block5_pool (MaxPooling2D) | (None, 4, 4, 512)    | 0       |
| Conv1 (Conv2D)             | (None, 1, 1, 64)     | 524352  |
| Conv2 (Conv2D)             | (None, 1, 1, 128)    | 73856   |
| Flatten (Flatten)          | (None, 128)          | 0       |
| Output (Dense)             | (None, 16)           | 2064    |

  
Total params: 15,314,960  
Trainable params: 600,272  
Non-trainable params: 14,714,688
```

```
In [ ]: plot_model(model2)
```







```
In [ ]: history = model2.fit(train_images_generator,validation_data=valid_images_generator,epochs=3,steps_per_epoch=steps,callbacks=callbacks_list)

Epoch 1/3
76/76 [=====] - 3298s 44s/step - loss: 2.1096 - accuracy: 0.3419 - val_loss: 1.7485 - val_accuracy: 0.
4548
Epoch 2/3
76/76 [=====] - 3266s 43s/step - loss: 1.6078 - accuracy: 0.4972 - val_loss: 1.5021 - val_accuracy: 0.
5415
Epoch 3/3
76/76 [=====] - 3274s 43s/step - loss: 1.4304 - accuracy: 0.5565 - val_loss: 1.4017 - val_accuracy: 0.
5696

In [ ]: # Save the weights using the `checkpoint_path` format
checkpoint_path = "/content/"
# Save the weights
model2.save_weights('./checkpoints2/my_checkpoint')

In [ ]: %tensorboard --logdir {log_dir}

Reusing TensorBoard on port 6006 (pid 811), started 2:45:22 ago. (Use '!kill 811' to kill it.)

In [ ]: vgg16base_model3 = tf.keras.applications.VGG16(weights='imagenet', include_top=False,input_shape=(*target_size, 3))
for layer in vgg16base_model3.layers[:-6]:
    layer.trainable = False
conv_layer1 = tf.keras.layers.Conv2D(512, (3, 3), activation='relu',padding='same',name='Conv1')(vgg16base_model3.output)
conv_layer1 = tf.keras.layers.Conv2D(64,4,4, activation='relu',padding='same',name='Conv1')(vgg16base_model3.output)
conv_layer2 = tf.keras.layers.Conv2D(128, (3, 3), activation='relu',padding='same',name='Conv2')(conv_layer1)
flatten = tf.keras.layers.Flatten(name='Flatten')(conv_layer2)
output = tf.keras.layers.Dense(16, activation='softmax',name='Output')(flatten)
model3 = tf.keras.models.Model(inputs=vgg16base_model3.inputs,outputs=output,name='Model_3')
model3.compile(optimizer='Adam', loss='categorical_crossentropy',metrics=['accuracy'])

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kerne
ls_notop.h5
58892288/58889256 [=====] - 0s 0us/step
58900480/58889256 [=====] - 0s 0us/step

In [ ]: model3.summary()

Model: "Model_3"

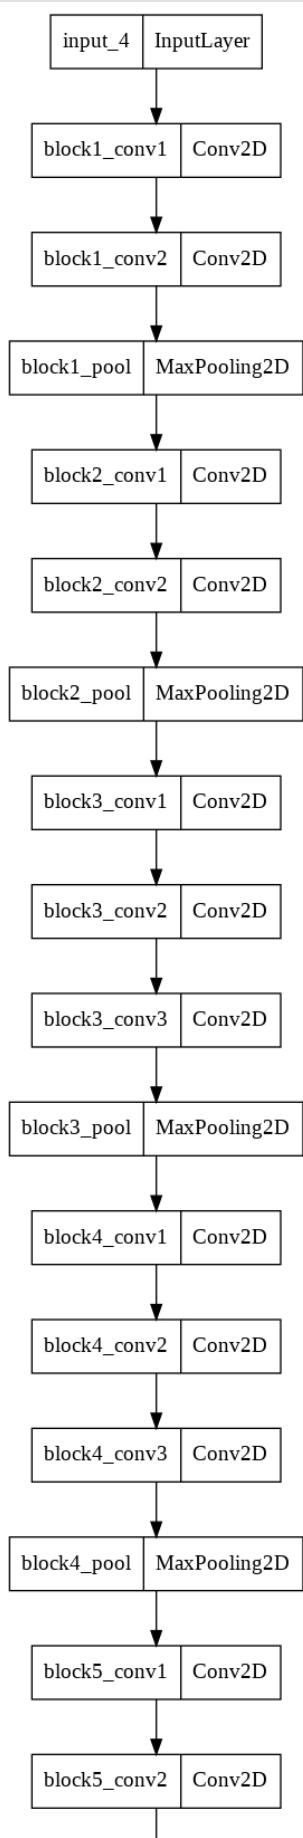
```

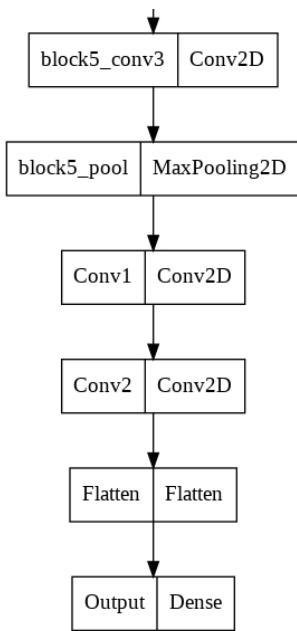
Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[None, 128, 128, 3]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
Conv1 (Conv2D)	(None, 1, 1, 64)	524352
Conv2 (Conv2D)	(None, 1, 1, 128)	73856
Flatten (Flatten)	(None, 128)	0
Output (Dense)	(None, 16)	2064

```
=====
Total params: 15,314,960
Trainable params: 10,039,504
Non-trainable params: 5,275,456
```

```
In [ ]: plot_model(model3)
```

```
Out[ ]:
```





```
In [ ]: import datetime
batch_size = 800
steps = len(train_images_generator.labels) // batch_size
early_stopping = keras.callbacks.EarlyStopping(patience=2)
tensorboard= tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,write_graph=True,write_grads=True)
callbacks_list = [early_stopping,tensorboard]
log_dir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
history = model3.fit(train_images_generator,validation_data=valid_images_generator,epochs=3,steps_per_epoch=steps,callbacks=callbacks_list)

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
Epoch 1/3
48/48 [=====] - 3134s 66s/step - loss: 2.6222 - accuracy: 0.1421 - val_loss: 2.3290 - val_accuracy: 0.
2473
Epoch 2/3
48/48 [=====] - 3099s 65s/step - loss: 2.1020 - accuracy: 0.3328 - val_loss: 1.8509 - val_accuracy: 0.
4126
Epoch 3/3
48/48 [=====] - 3150s 66s/step - loss: 1.7419 - accuracy: 0.4430 - val_loss: 1.6255 - val_accuracy: 0.
4923
```

```
In [ ]: %tensorboard --logdir {log_dir}
%load_ext tensorboard
%tensorboard --logdir {log_dir}

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

Reusing TensorBoard on port 6006 (pid 326), started 0:00:26 ago. (Use '!kill 326' to kill it.)
```

```
In [ ]: # List all data in history
print(history.history.keys())
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

